(Quasi-)Newton methods

Contents

1 Introduction

This course is largely based on the chapters 6 and 8 in [1].

1.1 Newton method

Newton method is a method to find the zeros of a differentiable non-linear function g, x such that $g(x) = 0$, where $g : \mathbb{R}^n \to \mathbb{R}^n$. Given a starting point x_0, Newton method consists in iterating:

$$x_{k+1} = x_k - g'(x_k)^{-1} g(x_k)$$

where $g'(x)$ is the derivative of g at point x. Here $g'(x_k)^{-1}$ is matrix in $\mathbb{R}^{n \times n}$. So each iteration requires to solve a linear system.

Applying this method to the optimization problem:

$$\min_{x \in \mathbb{R}^n} f(x)$$

consists in setting $g(x) = \nabla f(x)$, i.e. looking for stationary points. The iterations read:

$$x_{k+1} = x_k - \nabla^2 f(x_k)^{-1} \nabla f(x_k) \ .$$

Newton method is particularly interesting as its convergence is quadratic locally around x^*, i.e.:

$$\|x_{k+1} - x^*\| \leq \gamma \|x_k - x^*\|^2, \gamma > 0 \ .$$

Theorem 1 (Convergence of Newton method). *Let $g : \mathbb{R}^n \to \mathbb{R}^n$ assumed twice differentiable $\mathcal{C}^2$, and $x^* \in \mathbb{R}^n$ an isolated zero of g ($g(x^*) = 0$). Let us assume that $g'(x^*)$ is invertible. Then there exists a closed ball $\mathcal{B}$ centered on x^*, such that for every $x_0 \in \mathcal{B}$, the sequence x_k obtained with Newton algorithm stays in $\mathcal{B}$ and converges towards x^*. Furthermore, there is a constant $\gamma > 0$, such that $\|x_{k+1} - x^*\| \leq \gamma \|x_k - x^*\|^2$.*

PROOF. g' is continuous and $g'(x^*)$ is invertible so there exists a ball $\mathcal{B}(x^*, r_0)$ in which g' stays invertible and $(g')^{-1}$ is uniformly bounded by a constant $m > 0$. Using a Taylor expansion with integral form of the remainder one has:

$$-g(x_k) = g(x^*) - g(x_k) = g'(x_k)(x^* - x_k) + \int_0^1 g''(x^* + t(x_k - x^*)) \cdot (x^* - x_k)^2 t dt \ .$$

At each iteration one iterates $x_{k+1} = x_k - g'(x_k)^{-1} g(x_k)$. This implies

$$x_{k+1} - x^* = g'(x_k)^{-1} \left(\int_0^1 g''(x^* + t(x_k - x^*)) \cdot (x^* - x_k)^2 t dt \right) \ .$$

As g'' is also uniformly bounded by M, one obtains

$$\|x_{k+1} - x^*\| \leq \frac{mM}{2} \|x^* - x_k\|^2 \ .$$

Let $r = \min(r_0, \frac{2}{mM})$. It is easy to show that for x_0 in $\mathcal{B}(x^*, r)$ the sequence remains in the ball. Let

$$e_k = \frac{mM}{2} \|x^* - x_k\| \ .$$

One has that $e_{k+1} \leq e_k^2$, so if e_0 is less than 1, the sequence converges towards x^*.

As it is made clear by the proof, the convergence is guaranteed only if x_0 is sufficiently close to x^*. The method may diverge if the initial point is too far from x^* or if the Hessian is not positive definite. In order to address this issue of local convergence, Newton method can be combined with a line search method in the direction $d_k = -\nabla^2 f(x_k)^{-1} \nabla f(x_k)$.

Exercise 1. *Show that for a quadratic function*

$$f(x) = \frac{1}{2} x^\top A x - b^\top x + c, \ x \in \mathbb{R}^n$$

with A symmetric positive definite, Newton method converges in one iteration independently of the choice of x_0.

1.2 Variable metric methods

The idea behind variable metric methods consists in using iterations of the form

$$\begin{cases} d_k = -B_k g_k \ , \\ x_{k+1} = x_k + \rho_k d_k \ , \end{cases}$$

where $g_k = \nabla f(x_k)$ and B_k is a positive definite matrix. If $B_k = I_n$, it corresponds to gradient descent. Fixing $B_k = B$ leads to the following remark.

Remark. When minimizing

$$\min_{x \in \mathbb{R}^n} f(x)$$

one can set $x = Cy$ with C invertible (change of variable). Let us denote $\tilde{f}(y) = f(Cy)$. This leads to:

$$\nabla \tilde{f}(y) = C^\top \nabla f(Cy) \ .$$

Gradient descent applied to $\tilde{f}(y)$ reads:

$$y_{k+1} = y_k - \rho_k C^\top \nabla f(C y_k)$$

which is equivalent to

$$x_{k+1} = x_k - \rho_k C C^\top \nabla f(x_k) \ .$$

This amounts to using $B = CC^\top$. In the case where f is a quadratic form, it is straightforward to observe that this will improve convergence.

Theorem 2. *Let $f(x)$ a positive definite quadratic form and B a positive definite matrix. The preconditioned gradient algorithm:*

$$\begin{cases} x_0 = \text{fixed}, \\ x_{k+1} = x_k - \rho_k B g_k, \quad \rho_k \text{ optimal} \end{cases}$$

has a linear convergence:

$$\|x_{k+1} - x^*\| \le \gamma \|x_k - x^*\|$$

where:

$$\gamma = \frac{\chi(BA) - 1}{\chi(BA) + 1} < 1 \ .$$

Here $\chi(M) = \lambda_1/\lambda_n$, *i.e.,* it is the ratio between largest and lowest eigenvalues. It is the conditioning for the Euclidian norm.

Remark. The lower the conditioning of BA, the faster is the algorithm. One cannot set $B = A^{-1}$ as it would imply having already solved the problem, but this however suggests to use B so that it approximates A^{-1}. This is the idea behind quasi-Newton methods.

2 Quasi-Newton methods

2.1 Quasi-Newton relation

A quasi-Newton method reads

$$\begin{cases} d_k = -B_k g_k \ , \\ x_{k+1} = x_k + \rho_k d_k \ , \end{cases}$$

or

$$\begin{cases} d_k = -H_k^{-1} g_k \ , \\ x_{k+1} = x_k + \rho_k d_k \ , \end{cases}$$

where B_k (resp. H_k) is a matrix which aims to approximate the inverse of the Hessian (resp. the Hessian) of f at x_k. The question is how to achieve this? One can start with $B_0 = I_n$, but then how to update B_k at every iteration? The idea is the following: by applying a Taylor expansion on the gradient, we know that at point x_k, the gradient and the Hessian are such that:

$$g_{k+1} = g_k + \nabla^2 f(x_k)(x_{k+1} - x_k) + \epsilon(x_{k+1} - x_k) \ .$$

If one assumes that the approximation is good enough one has:

$$g_{k+1} - g_k \approx \nabla^2 f(x_k)(x_{k+1} - x_k) \ ,$$

which leads to the quasi-Newton relation, a.k.a. the *secant condition*.

Definition 1. Two matrices B_{k+1} and H_{k+1} verify the quasi-Newton relation if:

$$H_{k+1}(x_{k+1} - x_k) = \nabla f(x_{k+1}) - \nabla f(x_k)$$

or

$$x_{k+1} - x_k = B_{k+1}(\nabla f(x_{k+1}) - \nabla f(x_k))$$

The follow up question is how to update B_k while making sure that it stays positive definite.

2.2 Update formula of Hessian

The update strategy at iteration

$$\begin{cases} d_k = -B_k g_k \ , \\ x_{k+1} = x_k + \rho_k d_k \ , \end{cases}$$

is to correct B_k with a symmetric matrix Δ_k:

$$B_{k+1} = B_k + \Delta_k$$

such that quasi-Newton relation holds:

$$x_{k+1} - x_k = B_{k+1}(g_{k+1} - g_k)$$

with B_{k+1} positive definite, assuming B_k is positive definite. For simplicity we will write $B_{k+1} > 0$. We will now see different formulas to define Δ_k, under some assumptions that its rank is 1 or 2. We will speak about rank 1 or rank 2 corrections.

2.3 Broyden formula

Broyden formula is a rank 1 correction. Let us write

$$H_{k+1} = H_k + \sigma v v^\top \quad , \quad \sigma = \pm 1, v \in \mathbb{R}^n$$

The matrix H_{k+1} should verify the secant condition: $y_k = H_{k+1}s_k$, where $y_k = g_{k+1} - g_k$ and $s_k = x_{k+1} - x_k$. It follows that:

$$y_k = H_k s_k + (\sigma v^\top s_k)v \Rightarrow \exists \delta \in \mathbb{R}, v = \delta(y_k - H_k s_k)$$

Using the equality it leads to:

$$y_k - H_k s_k = \sigma \delta^2 [s_k^\top(y_k - H_k s_k)](y_k - H_k s_k)$$

this imposes that:

$$\sigma = \text{sign}[s_k^\top(y_k - H_k s_k)] \quad \delta = \pm |s_k^\top(y_k - H_k s_k)|^{-1/2}$$

This leads to:

$$H_{k+1} = H_k + \frac{(y_k - H_k s_k)(y_k - H_k s_k)^\top}{(y_k - H_k s_k)^\top s_k}$$

Using the matrix inversion lemma (Woordbury-Sherman-Morrison) it leads to:

$$B_{k+1} = B_k + \frac{(s_k - B_k y_k)(s_k - B_k y_k)^\top}{(s_k - B_k y_k)^\top y_k} \quad ,$$

also known as Broyden or SR1 formula.

Theorem 3. *Let f a quadratic form positive definite. Let us consider the method that, starting for x_0, iterates:*

$$x_{k+1} = x_k + s_k \quad ,$$

where the vectors s_k are linearly independent. Then the sequence of matrices starting by B_0 and defined as:

$$B_{k+1} = B_k + \frac{(s_k - B_k y_k)(s_k - B_k y_k)^\top}{(s_k - B_k y_k)^\top y_k} \quad ,$$

where $y_k = \nabla f(x_{k+1}) - \nabla f(x_k)$, converges in less than n iterations towards A^{-1}, the inverse of the Hessian of f.

PROOF. Since we consider here a quadratic function, the Hessian is constant and equal to A. This leads to:

$$y_i = \nabla f(x_{i+1}) - \nabla f(x_i) = A(x_{i+1} - x_i) = As_i, \forall i.$$

We have seen that B_{k+1} is constructed such that:

$$B_{k+1}y_k = s_k \quad .$$

Let us show that:

$$B_{k+1}y_i = s_i , i = 0, \dots, k-1 \quad .$$

By recurrence, let us assume that it is true for B_k:

$$B_k y_i = s_i , i = 0, \dots, k-2 \quad .$$

Let $i \leq k - 2$. One has

$$B_{k+1} y_i = B_k y_i + \frac{(s_k - B_k y_k)(s_k^\top y_i - (B_k y_k)^\top y_i)}{(s_k - B_k y_k)^\top y_k} \quad . \tag{1}$$

By hypothesis, one has $B_k y_i = s_i$ which implies that $y_k^\top B_k y_i = y_k^\top s_i$, but since $As_j = y_j$ for all j, it leads to:

$$y_k^\top s_i = s_k^\top A s_i = s_k^\top y_i \quad ,$$

The numerator in (1) is therefore 0, which leads to: $B_{k+1} y_i = B_k y_i = s_i$. One therefore has:

$$B_{k+1} y_i = s_i, \forall i = 0, \ldots, k \quad .$$

After n iterations, one has

$$B_n y_i = s_i, \forall i = 0, \ldots, n-1 \quad .$$

but since $y_i = As_i$, this last equation is equivalent to:

$$B_n A s_i = s_i, \forall i = 0, \ldots, n-1 \quad .$$

As the s_i form a basis of $\mathbb{R}^n$ this implies that $B_n A = I_n$ or $B_n = A^{-1}$. ∎

The issue with Broyden's formula is that there is no guarantee that the matrices B_k are positive definite, even if the function f is quadratic and $B_0 = I_n$ (σ can be -1). It is nevertheless interesting to have $B_n = A^{-1}$.

2.4 Davidon, Fletcher and Powell formula

The formula from Davidon, Fletcher and Powell is a rank 2 correction. It reads:

$$B_{k+1} = B_k + \frac{s_k s_k^\top}{s_k^\top y_k} - \frac{B_k y_k y_k^\top B_k}{y_k^\top B_k y_k} \quad . \tag{2}$$

The following theorem states that under certain conditions, the formula guarantees to have B_k positive definite.

Theorem 4. *Let us consider the update*

$$\begin{cases} d_k = -B_k g_k \quad , \\ x_{k+1} = x_k + \rho_k B_k g_k, \quad \rho_k \text{ optimal} \end{cases}$$

where $B_0 > 0$ is provided as well as x_0. Then the matrices B_k defined as in (2) are positive definite for all $k > 0$.

PROOF. Let $x \in \mathbb{R}^n \setminus \{0\}$. One has:

$$\begin{aligned} x^\top B_{k+1} x &= x^\top B_k x + \frac{(s_k^\top x)^2}{s_k^\top y_k} - \frac{(y_k^\top B_k x)^2}{y_k^\top B_k y_k} \\ &= \frac{y_k^\top B_k y_k x^\top B_k x - (y_k^\top B_k x)^2}{y_k^\top B_k y_k} + \frac{(s_k^\top x)^2}{s_k^\top y_k} \end{aligned} \tag{3}$$

If one defines the dot product $\langle x, y \rangle$ as $x^\top B_k y$ the equation above reads:

$$x^\top B_{k+1} x = \frac{\langle y_k, y_k \rangle \langle x, x \rangle - \langle y_k, x \rangle^2}{\langle y_k, y_k \rangle} + \frac{(s_k^\top x)^2}{s_k^\top y_k} \quad .$$

The first term is positive by Cauchy-Schwartz inequality. Regarding the second term, as the step size is optimal one has:

$$g_{k+1}^\top d_k = 0,$$

which implies

$$s_k^\top y_k = -\rho_k (g_{k+1} - g_k)^\top d_k = \rho_k g_k^\top B_k g_k > 0,$$

and so $x^\top B_{k+1} x \geq 0$. Both terms being positive, the sum is zero only if both terms are 0. This implies that $x = \lambda y_k$ with $\lambda \neq 0$. In this case the second term cannot be zero as $s_k^\top x = \lambda s_k^\top y_k$. This implies that $B_{k+1} > 0$. ∎

Remark. In the proof we used the fact that an optimal step size is used. The result still holds if one uses an approximate line search strategy for example using Wolfe and Powell's rule. In this case the point x_{k+1} is such that:

$$\phi'(\rho_k) = d_{k+1}^\top d_k \geq m_2 \|d_k\|^2, 0 < m_2 < 1,$$

which guarantees:

$$g_{k+1}^\top \frac{x_{k+1} - x_k}{\rho_k} > g_k^\top \frac{x_{k+1} - x_k}{\rho_k},$$

and therefore $(g_{k+1} - g_k)^\top (x_{k+1} - x_k) > 0$.

2.5 Davidon-Fletcher-Powell algorithm

One can now use the later formula in algorithm 1.

Algorithm 1: Davidon-Fletcher-Powell algorithm

Require: $\varepsilon > 0$ (tolerance), K (maximum number of iterations)
1:   $x_0 \in \mathbb{R}^n$, $B_0 > 0$ (for example I_n)
2:   **for** $k = 0$ to K **do**
3:     **if** $\|g_k\| < \varepsilon$ **then**
4:       break
5:     **end if**
6:     $d_k = -B_k \nabla f(x_k)$
7:     Compute optimal step size ρ_k
8:     $x_{k+1} = x_k + \rho_k d_k$
9:     $s_k = \rho_k d_k$
10:    $y_k = g_{k+1} - g_k$
11:    $B_{k+1} = B_k + \frac{s_k s_k^\top}{s_k^\top y_k} - \frac{B_k y_k y_k^\top B_k}{y_k^\top B_k y_k}$
12: **end for**
13: **return** x_{k+1}

This algorithm has a remarkable property when the function f is quadratic.

Theorem 5. *When f is a quadratic form, the algorithm 1 generates a sequence of directions $s_0, \ldots, s_k$ which verify:*

$$\begin{aligned} s_i A^\top s_j &= 0, \quad 0 \leq i < j \leq k, \\ B_{k+1} A s_i &= s_i, \qquad 0 \leq i \leq k. \end{aligned} \tag{4}$$

This theorem says that in the quadratic case, the algorithm 1 is like a conjugate gradient method, which therefore converges in at most n iterations. One can also notice that for $k = n - 1$ the equalities

$$B_n A s_i = s_i, \quad i = 0, \ldots, n-1,$$

and the fact that all s_i are linearly independent implies that $B_n = A^{-1}$.

Remark. One can show that in the general case (non-quadratic), if the direction d_k is reinitialized to $-g_k$ periodically, this algorithm converges to a local minimum $\hat{x}$ of f and that:

$$\lim_{k \to \infty} B_k = \nabla^2 f(\hat{x})^{-1} \ .$$

This implies that close to the optimum, if the line search is exact, the method behaves like a Newton method. This justifies the use of $\rho_k = 1$ when using approximate line search.

2.6 Broyden-Fletcher-Goldfarb-Shanno (BFGS) algorithm

The BFGS formula is a correction formula of rank 2 which is derived from the formula of DFP by swapping the roles of s_k and y_k. BFGS is named for the four people who independently discovered it in 1970: Broyden, Fletcher, Goldfarb and Shanno. The formula obtained allows to maintain an approximation H_k

of the Hessian which satisfies the same properties: $H_{k+1} > 0$ if $H_k > 0$ and satisfying the quasi-Newton relation:

$$y_k = H_{k+1}s_k.$$

The formula therefore reads:

$$H_{k+1} = H_k + \frac{y_k y_k^\top}{y_k^\top s_k} - \frac{H_k s_k s_k^\top H_k}{s_k^\top H_k s_k} \quad .$$

The algorithm is detailed in algorithm 2.

Algorithm 2: Broyden-Davidon-Goldfarb-Shanno (BFGS) algorithm

---

Require: $\varepsilon > 0$ (tolerance), K (maximum number of iterations)
1: $x_0 \in \mathbb{R}^n$, $H_0 > 0$ (for example I_n)
2: **for** $k = 0$ to K **do**
3:    **if** $\|g_k\| < \varepsilon$ **then**
4:       break
5:    **end if**
6:    $d_k = -H_k^{-1}\nabla f(x_k)$
7:    Compute optimal step size ρ_k
8:    $x_{k+1} = x_k + \rho_k d_k$
9:    $s_k = \rho_k d_k$
10:   $y_k = g_{k+1} - g_k$
11:   $H_{k+1} = H_k + \frac{y_k y_k^\top}{y_k^\top s_k} - \frac{H_k s_k s_k^\top H_k}{s_k^\top H_k s_k}$
12: **end for**
13: **return** x_{k+1}

Note that the direction d_k is obtained by solving a linear system. However in practice the update of H_k is done on Cholesky factorization of $H_k = C_k C_k^\top$ which implies that the complexity of BFGS is the same as DFP. The use of a Cholesky factorization is useful to check that H_k stays numerically positive definite, as this property can be lost due to numerical errors.

Remark. The BFGS algorithm has the same property as the DFP method: in the quadratic case it produces conjugate directions, converges in less than n iterations and $H_n = A$. Compared to DFP, BFGS convergence speed is much less sensitive to the use of approximate step size. It is therefore a particularly good candidate when combined with Wolfe and Powell's or Goldstein's rule.

Remark. The BFGS method is available in scipy as `scipy.optimize.fmin_bfgs`.

2.7 Limited-memory BFGS (L-BFGS) algorithm

The limited-memory BFGS (L-BFGS) algorithm is a variant of the BFGS algorithm that limits memory usage. While BFGS requires to store in memory a matrix of the size of the Hessian, $n \times n$, which can be prohibitive in applications such as computer vision or machine learning, the L-BFGS algorithm only stores a few vectors that are used to approximate the matrix H_k^{-1}. As a consequence the memory usage is linear in the dimension of the problem.

The L-BFGS is an algorithm of the quasi-Newton family with $d_k = -B_k \nabla f(x_k)$. The difference is in the computation of the product between B_k and $\nabla f(x_k)$. The idea is to keep in memory the last few low rank corrections, more specifically the last m updates $s_k = x_{k+1} - x_k$ and $y_k = g_{k+1} - g_k$. Often in practice $m < 10$, yet it might be necessary to modify this parameter on specific problems to speed up convergence. We denote by $\mu_k = \frac{1}{y_k^\top s_k}$. The algorithm to compute the descent direction is given in algorithm 3.

Commonly, the inverse Hessian B_k^0 is represented as a diagonal matrix, so that initially setting z requires only an element-by-element multiplication. B_k^0 can change at each iteration but has however to be positive definite.

Like BFGS, L-BFGS does not need exact line search to converge. In machine learning, it is almost always the best approach to solve ℓ_2 regularized Logistic regression and conditional random fields (CRF).

Algorithm 3: Direction finding in L-BFGS algorithm

Require: m (memory size)

1: $q = g_k$
2: **for** $i = k - 1$ to $k - m$ **do**
3:     $\alpha_i = \mu_i s_i^\top q$
4:     $q = q - \alpha_i y_i$
5: **end for**
6: $z = B_k^0 q$
7: **for** $i = k - m$ to $k - 1$ **do**
8:     $\beta = \mu_i y_i^\top z$
9:     $z = z + s_i(\alpha_i - \beta)$
10: **end for**
11: $d_k = -z$

Remark. L-BFGS is for smooth unconstrained problem. Yet, it can be extended to handle simple box constraints (a.k.a. bound constraints) on variables; that is, constraints of the form $l_i \leq x_i \leq u_i$ where l_i and u_i are per-variable constant lower and upper bounds, respectively (for each x_i, either or both bounds may be omitted). This algorithm is called $L - BFGS - B$.

Remark. The L-BFGS-B algorithm is available in scipy as `scipy.optimize.fmin_l_bfgs_b`.

3 Methods specific to least squares

3.1 Gauss-Newton method

When considering least square problems the function to minimize reads:

$$f(x) = \frac{1}{2} \sum_{i=1}^{m} f_i(x)^2 \ .$$

Newton method can be applied to the minimization of f. The gradient and the Hessian matrix read in this particular case:

$$\nabla f(x) = \sum_{i=1}^{m} f_i(x) \nabla f_i(x) \ ,$$

and

$$\nabla^2 f(x) = \sum_{i=1}^{m} \nabla f_i(x) \nabla f_i(x)^\top + \sum_{i=1}^{m} f_i(x) \nabla^2 f_i(x) \ .$$

If we are close to the optimum, where the $f_i(x)$ are assumed to be small the second term can be ignored. The matrix obtained reads:

$$H(x) = \sum_{i=1}^{m} \nabla f_i(x) \nabla f_i(x)^\top \ .$$

This matrix is always positive. Furthermore when m is much larger than n, this matrix is often positive definite. The Gauss-Newton method consists in using $H(x)$ in place of the Hessian. The method reads:

$$\begin{cases} x_0 = \text{fixed}, \\ H_k = \sum_{i=1}^{m} \nabla f_i(x_k) \nabla f_i(x_k)^\top, \\ x_{k+1} = x_k - H_k^{-1} \nabla f(x_k) \ . \end{cases}$$

3.2 Levenberg-Marquardt method

In order to guarantee the convergence of the Gauss-Newton method, it can be combined with a line search procedure:

$$x_{k+1} = x_k - \rho_k H_k^{-1} \nabla f(x_k) \ .$$

However in order to guarantee that the H_k stay positive definite a modified method, known as Levenberg-Marquardt, is often used. The idea is simply to replace H_k by $H_k + \lambda I_n$. One can notice that if λ is large, this method is equivalent to a simple gradient method. The Levenberg-Marquardt method reads:

$$
\begin{cases}
x_0 = \text{fixed}, \\
H_k = \displaystyle\sum_{i=1}^{m} \nabla f_i(x_k) \nabla f_i(x_k)^\top, \\
d_k = -(H_k + \lambda I_n)^{-1} \nabla f(x_k) \\
x_{k+1} = x_k + \rho_k d_k \ \ .
\end{cases}
$$

Remark. The Levenberg-Marquardt method is available in scipy as `scipy.optimize.leastsq`.

References

[1] J. Nocedal and S. J. Wright. *Numerical optimization.* Springer Series in Operations Research and Financial Engineering. Springer, New York, second edition, 2006.